

KnowledgeHub AI Sample Document

This PDF is intended for automated tests of the indexing pipeline.

It contains enough text to exercise parsing, cleaning, chunking, and embedding generation.

FastAPI is used to build APIs. Docker packages applications into containers. Retrieval-Augmented Generation (RAG) combines vector search with large language models. Embeddings convert text into numerical vectors that preserve semantic meaning. Chunking breaks large documents into smaller overlapping sections before embeddings are generated. ChromaDB stores embeddings for efficient semantic retrieval.